

THE ROAD TO A LIVE PRODUCTION TEST

Everything between the current build and running the agent against real data — grouped by dependency, honest about what's built and what isn't. Discovery can be tested soon; outreach is gated behind compliance.

- ✓ Done — already built & tested
- You — Marcus (infra / keys / decisions)
- Claude — remaining build work
- Gate — blocks the next stage

What "production testing" means here — two levels

Be clear which test you're preparing for; they have very different readiness bars.

- Discovery

Deploy + backfill, then browse a real, evidenced corpus of horror/sci-fi financiers in the dashboard. Reachable after Sections 1–3.
- The product

A *qualified, deduplicated, ranked, contactable* list — the actual deliverable. Needs Sections 4–5 (the Claude classifier + contacts).
- Outreach

Actually emailing/calling anyone. **Hard-gated** on Section 6 compliance — do not test outreach before it passes.

00 Where the build is now

Committed on the branch, PR #1, 20 tests green in CI.

✓	Schema with integrity constraints — budgets require provenance, claims require evidence, contacts verified-or-hidden Enforced by the database, not convention	Done
✓	ToS source gate — IMDb / IMDbPro / LinkedIn / Box Office Mojo physically uninsertable DB trigger + application check, seeded from the docs/02 verdicts	Done
✓	Wikidata + SEC EDGAR Form C ingestion — idempotent, evidence on every claim Fixture-tested; live response shapes still to be confirmed (see 3)	Done
✓	Dashboard skeleton with fail-closed auth gate + genre query Serves 503 in production if no password is set	Done
✓	CI (GitHub Actions) running the full suite on every push Against a real pgvector Postgres container	Done

01 Infrastructure & deploy

~20 minutes of clicking. Full steps in docs/11-deployment.md.

- ☐ Supabase project created; **vector** extension enabled You
Pro tier before real data — free tier pauses after a week idle, no backups
- ☐ Session-pooler connection string captured (kept out of chat) You
- ☐ Railway **worker** service — pre-deploy `npm run migrate`, start `npm run worker` You
- ☐ Railway **dashboard** service — build + start commands, public domain generated You
- ☐ Env vars set — `DATABASE_URL`, strong `DASHBOARD_PASSWORD`,
`NPM_CONFIG_PRODUCTION=false` You
- ☐ Migrations applied in prod (auto on worker deploy) & dashboard reachable, auth gate confirmed You

02 First data backfill

Runs on Railway, where network egress exists (this build container has none).

- ☐ Run Wikidata ingest — fills horror/sci-fi films + production companies (US/UK/EU/CA, 2016+) You
- ☐ Run SEC EDGAR ingest — fills issuer + funding-portal entities You
- ☐ Capture the stats JSON each run prints; hand it back Gate
Live EDGAR/SPARQL shapes were fixture-verified — first real run is where drift shows
- ☐ Fix any response-shape drift the stats reveal (high `fetchErrors` / zero rows) Claude
- ☐ Confirm the dashboard shows real financiers under the genre chips You

03 Core intelligence — qualification (Phase 2)

The difference between "a raw list" and "the product." Blocked on one key.

- ☐ ANTHROPIC_API_KEY set on the worker service (never pasted in chat) Gate
- ☐ Money-vs-craft classifier — separates financiers from creative credits, with confidence Claude
- ☐ Entity resolution — dedupe funds / SPVs across name variants and shell LLCs Claude
- ☐ Cluster rule + bucketing — **qualified_sub10m** vs mixed-scale / insufficient-data Claude
- ☐ Labeled eval set (50–100 entities) + precision floor for the qualified bucket
Precision prioritized over recall — a wrong "qualified" is worse than a miss Claude
- ☐ Ranking — genre affinity + warm signal (deal recency) + contactability Claude

04 Breadth & contacts (Phases 3–4)

Needed for a contactable list. Email + phone, verified per channel.

- ☐ Breadth sources — fund/prodco own sites, soft-money bodies, sales agents / MG distributors Claude
- ☐ Contact extraction — email *and* phone from entities' own public channels only Claude
- ☐ Verification — email-validation API + phone line-type lookup (keys needed) You + Claude
- ☐ Suppression / do-not-contact list operational across all channels Claude
- ☐ Assisted-draft view — personalized message, human sends from their own inbox Claude

05 Legal & compliance — outreach gate

Discovery testing needs none of this. Testing any outreach needs all of it.

- ☐ GDPR legitimate-interest basis recorded per stored contact (UK/EU in scope) Gate
- ☐ Opt-out / erasure honored across channels; suppression enforced before any surfacing Gate
- ☐ Phone rules — human-dial only (no autodialer/SMS), UK TPS/CTPS screening, business-line bias Gate
- ☐ CAN-SPAM elements in the draft template — identification, physical address, opt-out Gate
- ☐ Coverage caveat travels into every export ("discoverable, not complete") Claude
- ☐ Re-check ToS of any needs-license source before flipping it on You + Claude

06 Security & operations

- ☐ Replace any test `DASHBOARD_PASSWORD` with a strong secret; rotate if ever shared You
- ☐ Supabase backups on (Pro tier); secrets only in Railway env, never in repo/chat You
- ☐ Error logging + queue-health visibility; rate-limit politeness confirmed on live sources Claude
- ☐ Cost guardrails — Haiku/Opus routing, prompt caching, a spend alert on the Claude key You + Claude

07 Sign-off before calling it "ready to test"

- ☐ CI green on the PR; production build passes Claude
- ☐ End-to-end: ingest → classify → resolve → qualify → a ranked, evidenced list in the dashboard Claude
- ☐ Spot-check 5–10 entities against their evidence links — every claim traces to a source You + Claude
- ☐ False-positive control verified — a real sub-\$10M horror financier ranks; a mega-budget one-off is demoted Claude
- ☐ Define what the test measures — list precision, contactability rate, and a pass bar You

*Reflects the build at the time of writing (Phase 0 + Phase 1 ingestion complete, 20 tests green).
"Done" items are committed and tested; everything else is outstanding. Cost and vendor figures
elsewhere are planning-grade pending real quotes. The agent surfaces publicly discoverable
financiers — never a guarantee of completeness.*